

Urbit ~

The technical problem

- the high-level deterministic computer
 - official definition
 - entire lifecycle defined by a single frozen function
 - lifecycle semantics defined at the programmer level
 - existing approximations
 - machine VM: inherently low-level
 - JS, JVM, etc: transient, and not quite frozen
 - Lisp, Smalltalk machines: no functional definition
 - user experience: integrated OS/interpreter/DB
 - when a deterministic computer hits an undecidable problem?
- kelvin versioning
 - decreases by integers to absolute zero; Urbit is 5K

The human problem (1)

- the Internet as a client-server network has won
 - (HTTP = ATDT) \implies (FB = AOL)
- the Internet as a peer-to-peer network has failed
 - new wide-area protocols can no longer be introduced
 - even SMTP survives only by inertia
 - no one wants to self-host on a personal Linux server
 - Linux is layer 7 of the Internet
 - the Linux/Internet platform is unsalvageable as a P2P network

The human problem (2)

- the old platform is a fine substrate to layer over
 - on the client side, the browser already did this
- there's a plausible need for a new platform
 - we don't know that people don't want personal servers
 - we just know they don't want personal Unix servers on the Internet
- the obstacle to the personal server is admin cost
 - technical simplicity is a plausible therapy
- this “browser for the server side” is a clean slate
 - since we're layering over both the OS and the network

The one-function computer

- the simplest network: a global broadcast ethernet
 - routing is an optimization (content-centric networking)
 - packets are facts; the event log is a list of facts heard
- the simplest computer: a packet transceiver
- two ways to define a one-function computer:
 - lifecycle function: $L(\text{input history}) \rightarrow \text{resulting state}$
 - transition function: $T(\text{input event, state}) \rightarrow (\text{output actions, new state})$
- any practical L will converge to some T
 - output in a lifecycle function is an optional hint
 - a lifecycle function can define a boot sequence

Practical implementation

- event sourcing is popular these days
 - good tools for low-latency reliable logs (Kafka)
 - snapshot and append-only log is the normal DB design
 - every packet is a transaction, finalized when complete
- non-packet I/O can be event-sourced (libuv)
- decidability is a heuristic problem
 - interrupt a console; time out a packet
 - the log is an existence proof of computability
- nondeterminism feeds back into the event stream

Let's try it in Lisp

- it's easy to define a lifecycle function in Lisp
 - `(defun lifecycle (log) ((car log) (cdr log)))`
 - the first event is the function, the rest of history is the argument
 - of course we still have to write the function...
- now all we need is the one true perfect frozen Lisp
 - a lifecycle function makes extreme demands on interpreter precision
 - probably no “one-stage” interpreter can satisfy these demands
 - two stages: axiomatic untyped VM, user-level typed compiler
- lambda puts high-level features in the axioms
 - symbols, functions, variables, scopes, are all user-level features

The Urbit stack

- Nock: typeless, frozen interpreter
 - defined in 200 words, on a (readable) T-shirt
- Hoon: pure, strict, typed functional language
 - does not use category theory
 - compiles its own compiler to Nock
- Arvo: nonpreemptive operating system
 - written in Hoon
 - defines a transition function $T(\text{event}, \text{state}) \rightarrow (\text{actions}, \text{new state})$
 - defines a referentially transparent global namespace
 - interprets and sends untrusted, unreliable network packets
 - can update everything except Nock over the wire

Nock ideals

- a “functional assembly language”
- a Lisp without high-level affordances
 - symbols, variables, scope, syntax, etc
- no cyclic or infinite data structures
 - acyclic strict data is much easier to persist and transport
- efficient execution strategy
- fits on a T-shirt
- obviously perfect and will never need to change

Nock concepts

- a value in Nock is a *noun*
- a noun is an *atom* or a *cell*
 - an atom is an unsigned integer of any size
 - a cell is an ordered pair of any two nouns
 - cells are strict and acyclic and compare by value
- Nock is a function $*[\text{subject formula}] \rightarrow \text{product}$
 - the *subject* is the data; the *formula* is the function; the *product* is the result
 - any error produces nontermination, bottom, \perp

Nock spec (intrinsic)

?[a b]
?a

0
1

+ [a b]
+a

+ [a b]
1 + a

= [a a]
= [a b]
=a

0
1
=a

/[1 a]

a

/[2 a b]

a

/[3 a b]

b

/[(a + a) b]

/[2 /[a b]]

/[(a + a + 1) b]

/[3 /[a b]]

/a

/a

Nock spec (instructions)

[a [b c] d] [[a b c] *[a d]]

*[a 0 b] /[b a]

*[a 1 b] b

[a 2 b c] [[a b] *[a c]]

[a 3 b] ?[a b]

[a 4 b] +[a b]

[a 5 b] =[a b]

*[a 6 b c d] *[a 2 [0 1] 2 [1 c d] [1 0] 2 [1 2 3] [1 0] 4 4 b]

*[a 7 b c] *[a 2 b 1 c]

*[a 8 b c] *[a 7 [[7 [0 1] b] 0 1] c]

*[a 9 b c] *[a 7 c 2 [0 1] 0 b]

*[a 10 [b c] d] *[a 8 c 7 [0 3] d]

*[a 10 b c] *[a c]

*a

*a

Decrement in Nock and Hoon

```
[8
  [1 0]
  [ 8
    [ 1
      [ 6
        [5 [0 7] [4 0 6]]
        [0 6]
        [9
          2
          [[0 2] [4 0 6] [0 7]]
        ]
      ]
    ]
  ]
  [9 2 0 1]
]
```

```
:per    a=.
:pin    b=0
:loop
:if    =(a +(b))
      b
:moar(b +(b))
```

```
=>    a=.
=+    b=0
|-
?:    =(a +(b))
      b
$(b +(b))
```

Nock optimization

- $O(n)$ decrement is nifty, but not practical
- solution: a sufficiently smart interpreter
 - just recognize any decrement formula, and execute it efficiently
 - also add, multiply, and all other common intrinsics
- wait, we don't have to recognize *every* decrement
 - just the decrement in the standard library
- solution: hint-register and match important functions
 - these *jets* are like software device drivers
 - user code remains pure, but declares a semantic identity
 - an interpreter which recognizes this identity can optimize it

Hoon ideals

- compiler compiles itself from source to Nock
 - does *not* have to be frozen; anything above Nock must self-update
- pure, strict, higher-order typed functional language
- transformation to Nock is simple, like C to assembly
- requires no particular mathematical aptitude
 - and does not use category theory
- encourages lower-order, more imperative style
 - DSLs considered harmful
- still almost as expressive as Haskell

Hoon concepts

- Hoon is a typed macro assembler for Nock
 - type inference and code generation combined are 1.500 lines of Hoon
- Computes [subject expression] → product
 - “subject-oriented programming”
- the Hoon compiler is a parser and a generator
 - the parser (vast) computes source → expression
 - the generator (ut) computes [type expression] → [type formula]
 - its input is the subject type (domain) and the expression
 - its output is the product type (range) and the Nock formula
- Hoon infers only forward, without unification
 - a (slightly) less intelligent inference algorithm is a UI win
 - it's good to coerce the product of any entry point, just for doc reasons

Never say type

- Learn the secret language of Hoon
 - an expression (AST noun) is a twig
 - a type (as in set of nouns) is a span
 - a type (as in constructor/declaration) is a mold
 - a type (as in MIME type) is a mark
- Hoon is a pure prototype language
 - there is no syntax for a span; it is only defined as the range of a computation
- A mold is a normalizing function on an arbitrary noun
 - a true mold is idempotent, so `=((mold x) (mold (mold x)))`
 - the only time we actually *call* a mold is to validate network data

Basic span concepts

- twigs are boring once you understand spans
- mold for a slightly simplified span:

```
++ span
$@ $? $noun
      $void
== $% {$atom p/term q/(unit atom)}
      {$cell p/span q/span}
      {$core p/span q/(map term span)}
      {$face p/term q/span}
      {$fork p/(set span)}
      {$hold p/span q/twig}
==
```

- missing only: aliases, variance and typeclasses

Unboring spans: \$atom

- `{ $atom p / term q (unit atom) }`
 - if `q` is set, `u.q` is the only atom in the range (constant)
 - `p` is an aura, a symbol which describes units/presentation/constraint
 - auras are unenforced conventions
 - auras specialize by extension right
 - `@t` for UTF-8 text, `@ta` ASCII, `@tas` ASCII with symbol constraint
 - `@s` for signed integer, `@sx` for signed hexadecimal integer
 - auras can be cast upward or downward, but not across

Unboring spans: \$score

- `{ $score p/span q/(map term twig) }`
 - a core is a `[battery payload]` cell
 - `p` is the span of the payload
 - `q` is a table of computed attributes, or *arms*
 - the battery is the tree of the arm formulas
 - the subject of each arm is the core
 - single namespace searches battery first, then payload
- a core is the general case of functions and objects
 - a function in Hoon is a gate: a special case of core
 - a gate has one nameless arm (\$) and a payload `[sample context]`
 - a method in Hoon is an arm which produces a gate
 - an object in Hoon is a core whose arms are all methods

Advanced span theory

- the `$hold` form is manual laziness
 - but also lets conservative worklist algorithms prune recurrent traverses
- branch conditions are analyzed for type inference
 - this allows classic functional pattern matching
- two kinds of polymorphism: variance and genericity
 - polymorphism is about compatibility of mutated cores (Liskov substitution)
 - variance: does mutant payload match original payload?
 - collect all four: covariance, contravariance, invariance, bivarience
 - genericity: does battery formula work with mutant payload?
 - generic (wet) arms expand inline like macros

Syntax design

- Hoon syntax is twig syntax
- twig syntax is functionally complex and looks gnarly
 - but everyone who learns it is surprised at how easy it was
- Hoon solves three problems with functional syntax
 - expressions grow downward and to the right
 - solution: backstep indentation
 - either terminator piles or significant whitespace
 - solution: self-terminating form
 - hard to distinguish special forms from symbols
 - solution: no macros, marked keywords / runes

Twig structure

- a twig is the AST expression
 - which compiles to a Nock formula, which defines a function of the subject
- any cell [twig twig] is also a twig
 - which constructs the cell of its subtrees, like Lisp cons
 - Hoon is tuple-centric, not list-centric, because types work
- any other twig is a tagged union, [stem bulb]
 - the head of the twig is an atom, its stem
 - the stem is a 2-4 byte term (ASCII symbol)
 - the shape of the tail depends on the stem
 - but usually a tuple or list of twigs

Regular forms, flat and tall

- every stem defines its own bulb
 - most are 1-ary, 2-ary, 3-ary or 4-ary tuples; some are n-ary lists
 - every stem with a tuple/list bulb has a regular form
- regular forms come in two forms: *flat* and *tall*
 - a flat twig is delimited by parens, and separates subtwigs by one space:
`:if(a b c)`
 - a tall twig has no delimiters, and separates subtwigs by 2+ spaces:
`:if a
 b
 c`
- goal: look like a procedural expression-statement mix

Runes and irregular forms

- irregular form: always flat, always ASCII, and always a twig
- a regular form can use a keyword or a *rune*
 - keyword is colon-prefixed: :if(a b c)
 - rune is a digraph: ?:(a b c)
 - first character defines role (eg, all ? runes are conditionals)
- ASCII reloaded:

ace [1 space]	gal <	pal (
bar	gap [>1 space, nl]	par)
bas \	gar >	sel [
buc \$	hax #	sem ;
cab _	hep -	ser]
cen %	kel {	sig ~
col :	ker }	soq '
com ,	ket ^	tar *
doq "	lus +	tec `
dot .	pam &	tis =
fas /	pat @	wut ?
zap !		

Two forms of FizzBuzz

```
:gate end/atom
:var count 1
:loop
:cast (list tape)
:if =(end count)
    ~
:cons
    :if =(0 (mod count 15))
        "FizzBuzz"
    :if =(0 (mod count 5))
        "Fizz"
    :if =(0 (mod count 3))
        "Buzz"
    "{<count>}"
:moar(count (add 1 count))
```

```
|= end/atom
=/ count 1
|- ^- (list tape)
?: =(end count)
    ~
:_ $(count (add 1 count))
?: =(0 (mod count 15))
    "FizzBuzz"
?: =(0 (mod count 5))
    "Fizz"
?: =(0 (mod count 3))
    "Buzz"
"{<count>}"
```

Arvo concepts

- the Arvo kernel is a Hoon core
 - with a fixed battery exporting a few arms
 - including a transition function $T(\text{event}) \rightarrow (\text{actions}, \text{new core})$
 - the kernel ABI is frozen at the Nock level
 - so any event can replace Arvo or even Hoon
 - as long as it can build a core that looks the same to Nock
- Arvo proper is ~600 lines
 - internal event cascade with causal stack
 - global typed referentially transparent namespace
 - load and reload kernel modules (vanes), like baby kernels

Arvo vanes (kernel modules)

- %ames, encrypted packet network
- %behn, timers
- %clay, revision-controlled typed filesystem
- %dill, console
- %eyre, HTTP client/server with reactive apps
- %ford, functional build system
- %gall, applications

Urbit identity concepts

- “Urbit” just means the Nock/Hoon/Arvo stack
 - and more specifically the PKI / identity model
 - “an urbit” means one event history / state / instance/ node, with one identity
 - identity creation is part of the boot process
 - the identity is a ship, the instance is a *pier*
- one system solves:
 - human-memorable cryptographic identity
 - P2P packet routing address
 - base of global immutable namespace path
- assault on Zooko’s triangle
 - trivial solution: identity is 128-bit public-key hash
 - tradeoff: memorable, but not meaningful

A civil address space

- phonemic base-256 makes numbers memorable
 - 0x802a.136d in @ux; .128.42.19.109 in @if; ~patnub-tarlud in @p
- a 128-bit “wild” ship is called a comet
 - but shorter numbers are more memorable, hence more valuable
- a 64-bit “civil” address space can overload it
 - 64-bit ship: *moon*; ~padfes-sopden-difmyl-padtul; connected device
 - 32-bit ship: *planet*; -difmyl-padtul; human being
 - 16-bit ship: *star*; ~mocryg; minor infrastructure
 - 8-bit ship: *galaxy*; ~num; major infrastructure
- each tier is initially signed by its half-width parent
 - but signs its own key update; renewal is revocation
 - ships can be traded like bitcoin, but low-frequency “spends” don’t need a blockchain
 - galaxies are “premined” in the kernel source

On purpose and hindsight

- but was it really necessary to invent all this crap?
- a personal server is a social server
 - when two humans socialize, they should exchange messages directly
 - and not fall back into the degenerate case of a central server
- centralized programming is always easier
 - in today's infrastructure it's orders of magnitude easier
- the criterion: *difficulty of distributed programming*
- a true personal server must solve this problem
 - some impersonal servers wouldn't mind as well

Programmer experience

- dereference the global immutable typed namespace
- application state is permanent, no database required
- update source code propagates reactively via revision control
 - application type changes require typed state adapters
- two messaging patterns, poke and peer
 - no abstraction leakage versus local communication
 - %eyre lets web clients poke and peer over HTTP
- a poke is a typed, transactional message
 - exactly-once delivery (even though it's impossible)
 - no return data if transaction succeeds
 - message is automatically validated
 - backward-compatible protocol updates cause no errors in a live network
 - passed to the receiving program as a typed event
 - end-to-end acknowledgments mean single error mode
 - sender queues until message delivered or rejected
 - P2P network delivers anywhere
 - authenticated and encrypted, of course
- a peer creates a subscription which streams typed diffs

System status

- 30,000 lines of Hoon (Hoon, Arvo, all vanes, and basic apps)
- open source and patent-free
- urbit.org
 - served by Urbit (behind nginx :-)
- github.com/urbit
- somewhere between alpha and beta quality
 - small live network, “used in anger”, working shell and chat apps
 - Arvo internal interfaces mostly stable
 - global flag day (“continuity breach”) every few months
- documentation historically sucks, but getting better
 - Hoon documentation is now adequate
 - next phase is Arvo documentation
- ready for self-hosted address-space crowdsale